Drift-Sense: Predicting Position Drift from Inertial Sensor Data

Dadi Dinesh | Enrollment ID: 2401010136
Machine Learning Project
Code: https://github.com/Dadi-Dinesh/Drift-Sense

Abstract—Drift-Sense studies whether motion statistics from an inertial measurement unit can predict short-window position drift. The pipeline aligns two EuRoC Machine Hall flights with ground truth, builds two-second windows, constructs labels using strapdown integration and extracts 104 sensor features. Mean prediction, ordinary least squares, Ridge and Lasso are compared using leave-one-flight-out evaluation with blocked inner tuning. Committed results report a mean Lasso MAE of 0.0250 m and mean fold R^2 of 0.756, compared with baseline MAE of 0.0453 m. The 44.7% reduction measures drift-prediction error. The system does not yet correct trajectories. Only two flights support these development results, so unseen-flight evaluation remains essential.

I. PROBLEM AND OBJECTIVES

Problem statement

An IMU measures angular rate and specific force. An inertial navigation calculation uses angular rate to update orientation, rotates specific force into the world frame, accounts for gravity, and integrates acceleration to obtain velocity and position. Measurement biases and numerical errors accumulate during these updates. The resulting position estimate can differ substantially from a reference trajectory.

The project asks whether a window of raw IMU data contains enough information to estimate the magnitude of this accumulated error. This is a supervised regression task: the input consists of sensor-derived features and the target is a distance in metres. The model estimates the error of a particular short-window integration procedure.

Project objectives

Build consistent drift labels, extract interpretable sensor features, compare simple regression models, and assess generalization across flights without mixing overlapping windows.

An illustrative example

Suppose the true endpoint is (5, 3, 2) m and IMU integration produces (6, 3, 2) m. The label is a drift magnitude of 1 m. If the ML model estimates 1.03 m, its absolute prediction error is 0.03 m. This example separates the physical drift being predicted from the regression error used to assess the model.

Implemented scope

The implementation reads recorded gyroscope and accelerometer data. Camera images, a visual front end, live sensor ingestion and trajectory correction are outside the implemented main pipeline. Although the Python package is named vio_drift, the present experiment models IMU-only integration drift. The deep-learning and correction directories are placeholders.

II. STEP 1: DATA AND PREPARATION

The configured dataset consists of two easy Machine Hall flights from EuRoC MAV. Each flight supplies IMU readings and ground-truth state estimates. The committed quality summary records no missing values in the two raw streams and approximately 200 Hz sampling. [3, 6]

TABLE I

Property	MH 01 easy	MH 02 easy
Raw IMU rows	36,820	30,400
Ground truth rows	36,382	29,993
Recorded duration	184.095 s	151.995 s
Alignment rows dropped	431	400
Alignment drop rate	1.17%	1.32%
Two second windows	720	592

Loading and timestamp alignment

loader.py expects seven IMU columns and seventeen ground-truth columns. It assigns project column names, converts timestamps to integers, sorts by timestamp and removes duplicates. Ground-truth quaternions are normalized. align.py uses a nearest-time merge and drops matches farther than 20 ms from an IMU timestamp. The signed matching gap remains available for inspection.

Window configuration

The current configuration uses a two-second duration and a 0.25-second stride. A window contains approximately 400 readings across six channels. Neighboring windows share 87.5 percent of their duration. Window creation uses timestamps rather than a fixed number of rows, while an expected-row check rejects windows with an excessive sample-count difference.

The raw-window array has shape (N, 400, 6), with channels ordered gx, gy, gz, ax, ay, az. Slight sample-count differences are padded with the final sample or truncated. The statistical feature code subsequently assumes a nominal 200 Hz rate for FFT and jerk calculations.

III. SYSTEM ARCHITECTURE

Drift Sense separates reusable functions in src/vio_drift from command-line orchestration in scripts and explanatory experiments in notebooks. Files on disk form the interfaces between stages. The training path consumes statistical features and labels, while the physics-feature path supports separate diagnostics.

TABLE II

Stage	Implementation	Output
Load and align	data/loader.py and align.py	Aligned IMU and ground truth
Window and label	data/windowing.py and labels/	Window labels and raw arrays
Extract features	features/ and make_features.py	104 features per window
Fit and evaluate	models/ml/ and train_eval.py	Fold metrics and CV tables
Fit final model	train_final.py	Lasso pipeline and metadata

Training data contracts

windows_<sequence>.csv contains sequence and window identifiers, start and end timestamps, raw error magnitude and axis components, plus bias-corrected diagnostic labels. windows_raw_<sequence>.npy stores the raw sensor windows. features_<sequence>.csv stores identifiers, the 104 unscaled features and the primary target. splits.json records the outer flight assignments.

Training and inference

During training, ground truth creates the target labels. The deployed statistical model accepts the 104 sensor features without ground-truth columns. Its saved StandardScaler applies the same preprocessing used during fitting. Inference must preserve feature names, order, units, sampling assumptions and window duration.

The final artifact is models/ml/lasso_err_mag_m.pkl with a JSON metadata companion. The code provides a load-and-predict example, but a complete streaming inference service and input-quality checks are not yet implemented. The physics diagnostic has different assumptions because it uses ground-truth starting states and training-flight calibration biases.

IV. STEP 2: DRIFT LABEL CONSTRUCTION

Starting state and orientation

For each window, drift_label.py reads the first aligned ground-truth position p0, velocity v0 and quaternion q0. A quaternion represents three-dimensional orientation. The integrator starts from this reference state and then advances using the window's raw IMU measurements. Resetting every window isolates local integration error from drift accumulated before the window.

Integration equations

With the configured body-to-world rotation R(q), the acceleration update is a_world = R(q) f_body + g_world. The gravity vector is (0, 0, -9.81) m/s². For each timestamp interval Δt, position updates as p_next = p + v Δt + 0.5 a_world Δt², and velocity updates as v_next = v + a_world Δt. A rotation increment derived from angular rate advances the quaternion, which is normalized after multiplication.

The implementation compares the predicted endpoint with ground truth at the final retained row. The configured interval is half-open, so a regular 400-sample window at 200 Hz spans approximately 1.995 seconds between its first and last samples. The nominal two-second duration describes window selection.

Raw error and diagnostic error

The signed vector error is e = p_integrated - p_true. Its components become err_dx, err_dy and err_dz. The main regression label is err_mag_m = ‖e‖2. It is nonnegative and measured in metres.

A second integration subtracts the dataset's per-row gyroscope and accelerometer bias estimates. The resulting err_*_bc fields form an oracle diagnostic. The saved mean residual magnitude is 0.0631 m for MH_01_easy and 0.0540 m for MH_02_easy, substantially below the raw means near 1.21 m. This supports a strong bias contribution, but does not establish a precise additive percentage decomposition of vector-error magnitude. [3, 5]

V. STEP 3: FEATURE ENGINEERING

Each window becomes a vector of 104 numerical summaries. The feature families describe signal level, variation, frequency content and relationships between channels. Features use raw sensor readings, while ground-truth quantities remain outside the main statistical feature vector. [5, 7]

TABLE III

Family	Count	Examples and meaning
Time domain	60	Ten statistics on each of six channels
Frequency domain	18	Dominant frequency, energy and entropy per channel
Cross axis and magnitude	26	Vector magnitudes, jerk and fifteen correlations

Time domain statistics

For each channel, time_domain.py calculates mean, standard deviation, variance, RMS, minimum, maximum, peak-to-peak range, skewness, excess kurtosis and mean-centered zero-crossing rate. Centering zero crossings avoids letting a gravity offset dominate that feature. Constant channels receive zero skewness and kurtosis.

Frequency and cross axis statistics

frequency_domain.py removes the channel mean, computes a real FFT and summarizes dominant frequency, summed spectral energy and normalized spectral entropy. cross_axis.py computes acceleration and gyroscope magnitudes, finite-difference jerk at 200 Hz, and Pearson correlations for every channel pair. Undefined correlations for constant channels become zero.

Scaling and redundancy

Features have different units and scales. StandardScaler fits only on training rows inside each model pipeline, including inside cross-validation. Related features can be highly correlated, such as standard deviation and variance. Regularization reduces the instability that can arise when a linear model fits many correlated predictors to a small number of independent flights.

VI. STEP 4: MODELS AND EVALUATION

TABLE IV

Model	Training behavior
Mean baseline	Predicts the training label mean for every window
Ordinary least squares	Fits a linear combination of standardized features
Ridge regression	Adds an L2 penalty to shrink coefficient magnitudes
Lasso regression	Adds an L1 penalty and can set coefficients to zero

For a feature vector x, the linear prediction is y_hat = wTx + b. OLS minimizes squared error. Ridge adds $\alpha\|w\|2^2$ to the summed squared loss. scikit-learn Lasso uses (1/2n) times the summed squared loss plus $\alpha\|w\|1$. The alpha grids contain 17 Ridge values from 10^{-3} to 10^{5} and 16 Lasso values from 10^{-6} to 10^{-1}.

Outer flight split

The first fold trains on MH_02_easy and tests on MH_01_easy. The second reverses those roles. This prevents the heavily overlapping windows of one recording from appearing on both sides of the outer split. The number of independent flights remains two, even though the dataset contains 1,312 overlapping examples.

Inner tuning

Within each outer fold, the training flight is divided into five contiguous blocks. Each block serves as validation once, with the other blocks used for fitting. Sixteen nearby windows on either side are purged from training. At the current stride this is a four-second exclusion, exceeding the two-second window duration.

The implemented selection rule chooses the alpha with the lowest median block MAE. It reduces the influence of the unusually difficult early-flight block. The code also reports mean-error and one-standard-error selections. The median rule changes which operating segments influence tuning and therefore needs evaluation on independent flights.

Metrics and interpretation

MAE is the average absolute difference between predicted and true drift magnitude. RMSE gives larger residuals more weight. R^2 measures squared-error improvement relative to predicting the test set's own mean. Prediction bias is the mean of prediction minus target. Results average metrics equally across the two folds rather than pooling all windows.

Evaluation qualification

Diagnostic scripts compared window lengths and tuning rules using these same flights. Their use in design selection means the reported scores are development results rather than a final untouched benchmark. A new flight-level test set is needed after freezing preprocessing, model selection and tuning rules. [2, 8]

VII. RESULTS AND ERROR ANALYSIS

The committed results identify Lasso as the best of the four main models by mean held-out MAE. The table reports equal-weight averages over the two outer folds. The values describe prediction of local raw-IMU drift magnitude. [2]

TABLE V

Model	MAE cm	RMSE cm	Mean R^2
Mean baseline	4.53	6.91	-0.032
OLS	2.99	4.15	0.580
Ridge	2.70	3.49	0.717
Lasso	2.50	3.23	0.756

The relative MAE reduction is (0.0452505 - 0.0250447) / 0.0452505 ≈ 44.7 percent. Lasso's fold MAEs are approximately 2.59 cm on MH_01_easy and 2.42 cm on MH_02_easy. Its mean prediction bias is approximately -0.12 cm.

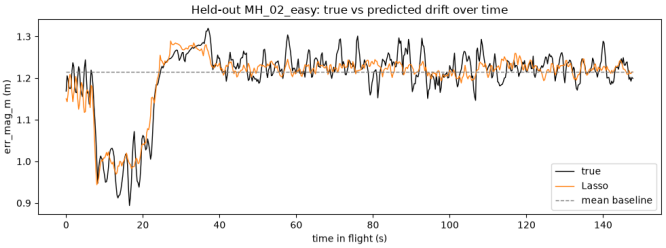


Fig. 1. Committed notebook output for MH_02_easy showing true and predicted drift over time. Source [8].

Generalization and flight segments

OLS has the largest gap between training and held-out MAE. Its held-out to training ratio is approximately 1.81, compared with 1.38 for Ridge and 1.24 for Lasso. These results support regularization within this limited experiment.

The report's segment analysis shows Lasso R^2 of 0.86 and 0.84 in the first fifth of the two flights, but only 0.10 and 0.17 in the remaining portions. The model follows the large early-flight drift change more successfully than rapid fluctuations during normal flight. A whole-flight R^2 of 0.756 should therefore be interpreted alongside segment-level performance.

VIII. STEP 5: FINAL MODEL

Saved model

train_final.py chooses alpha using leave-one-flight-out validation over both available flights, then refits Lasso on every window. The saved JSON records alpha = 0.0046415888, 1,312 windows, 104 input features and seven nonzero coefficients. Because the final model uses both flights for fitting, its training score is not an independent generalization measurement. [4]

TABLE VI

Retained feature	Interpretation
gx_max	Largest x-axis angular-rate reading
gy_mean	Mean y-axis angular rate
ax_rms	RMS of x-axis specific force
az_mean	Mean z-axis specific force
gz_dom_freq	Dominant frequency of z-axis angular rate
az_spec_entropy	Frequency spread of z-axis specific force
corr_ax_az	Correlation of x and z accelerometer channels

The saved pipeline retains the 104-feature input schema. Applications must supply the complete expected feature list in order, even though most Lasso coefficients are zero.

IX. LIMITATIONS AND CONCLUSION

Scientific limitations

Only two easy flights from a shared sensor setting support the current results. Overlapping windows are correlated, so the row count overstates the amount of independent evidence. The label measures a locally reset integrator and uses ground-truth initialization. Broader claims about continuous navigation, aggressive motion or a different IMU require separate experiments.

The model outputs a scalar magnitude, which alone cannot specify the direction of a position correction. A correction system would need signed error components or another directional estimate and must demonstrate improved trajectories over time. No such corrected-trajectory result is currently implemented.

Engineering improvements

First, preserve IDs consistently after label filtering and correct quaternion propagation in the alternate rotation branch. Second, update one-second documentation to the active two-second configuration. Third, strengthen semantic input validation and stop suppressing all training warnings. The hard-coded cross-validation gap should derive from window and stride settings when those become configurable.

TABLE VII

No.	Next work	Evidence required
1	Freeze pipeline and test new flights	MAE and R^2 per unseen flight and segment
2	Compare physics and nonlinear models	Same data splits and input assumptions
3	Model signed axis errors	Direction accuracy and residual analysis
4	Evaluate trajectory correction	Reduced continuous trajectory error
5	Explore temporal deep learning	Independent gains over simpler baselines

Conclusion

Drift Sense implements an interpretable supervised pipeline for predicting short-window IMU integration error. The saved experiments show that Lasso improves on a mean baseline and unregularized regression on the two available flights. The project's strongest demonstrated outcome is a reproducible feature and regression workflow, accompanied by useful diagnostics. Further validation should focus on untouched flights, realistic initialization and consistent evaluation before extending to live drift monitoring or correction.

X. REPRODUCIBILITY AND REFERENCES

The source snapshot is 5f24a96de9e67c965bcaa93c8524a013a37fbdf9. Reported model metrics come from committed experiment outputs. The complete raw-data pipeline was not independently rerun for this report. Source review parsed 38 nonempty Python files and directly executed 18 existing test functions.

Run the scripts in order: verify_data.py, make_windows.py, run_eda.py, make_features.py, train_eval.py and train_final.py. Run the repository test suite with python -m pytest tests/. Freeze configuration before evaluating new recordings.

[1] D. Dinesh, Drift-Sense source repository. https://github.com/Dadi-Dinesh/Drift-Sense. Snapshot dated 26 October 2026.

[2] Drift-Sense, results/metrics/err_mag_m.csv. Equal-weight fold metrics for the main regression comparison.

[3] Drift-Sense, results/eda/label_summary.csv and quality_summary.csv. Dataset counts and integration-label summaries.

[4] Drift-Sense, models/ml/lasso_err_mag_m.json and scripts/train_final.py. Final model settings and input schema.

[5] Drift-Sense, src/vio_drift/data, labels, features and models/ml. Processing, integration, feature and model implementation.

[6] EuRoC MAV dataset, as used and cited in the Drift-Sense repository. MH_01_easy and MH_02_easy sequences. https://doi.org/10.3929/ethz-b-000690084.

[7] Drift-Sense, notebooks/04_feature_engineering.ipynb. Feature exploration and feature selection.

[8] Drift-Sense, notebooks/05_model_training.ipynb, notebooks/06_model_evaluation.ipynb, scripts/tuning_check.py and reports/midsem/sections/results.tex. Evaluation and segment analysis.